

Benjamin Taoson 赖建铭

Senior AI Algorithm Engineer

LLM post-training · Agent systems · Research & delivery

benjamindaoson@gmail.com · github.com/Benjamindaoson

Profile

Over eight years in enterprise technology and delivery, working in AI since 2021 and LLM and agent systems since 2023. My work spans business analysis, architecture, core development, deployment, and iteration. Research interests include LLM post-training, VLA policies, and world models.

Experience

Jan 2023–Present

Senior AI Algorithm Engineer

Yaowu Technology (Shenzhen) Co., Ltd.

SANY global websites · Pre-sales solution agent

Led the architecture and delivery of an equipment-selection agent, connecting requirements, retrieval, business tools, and recoverable multi-turn tasks.

- **Runtime and delegation:** Built a DeepAgents / LangGraph supervisor with four specialist types. Routed direct calls or delegation by requirement completeness and dependencies, using isolated contexts, scoped tools, and structured outputs.
- **State and changing requirements:** Modeled six requirement groups in Pydantic. Delta merges and dependency invalidation enable local recomputation; offline equipment-model identification and disambiguation rose from over 80% to over 90%.
- **Retrieval and constraints:** Combined hard filters, BM25, Qwen3 embeddings, RRF, and Qwen3 reranking over roughly 40,000 records. Offline Top-5 candidate recall improved from about 75% to over 90%.
- **Governed tools and recovery:** Wrapped four tools for product queries, comparison, dealer matching, and CRM. Checkpoints, classified retries, business keys, and deduplication enable recovery without duplicate leads.
- **Trace evaluation and regression:** Evaluated over 100 multi-turn tasks with LangSmith / Langfuse traces, rule checks, and failure attribution. Success on a frozen regression set rose from over 70% to about 85%.

Jan 2021–Dec 2022

AI Engineer · NLP

Zhiguang Electric

Applied NLP to tender preparation and supported SAP integration across enterprise workflows.

- Used OCR, BERT / ERNIE, NER, and Elasticsearch / BM25 for tender parsing, requirement extraction, clause matching, compliance checks, and prior-material retrieval.
- Handled SAP requirements, data mapping, API integration, and testing across procurement, production, inventory, and finance.

Aug 2018–Dec 2020

Digital Solutions Engineer / Technical Project Manager

Zhiguang Electric

Delivered manufacturing ERP / CRM discovery, process design, integration, and launch within cross-functional projects involving over 30 people.

Selected projects

AI BA · Business Analytics

Built a supervisor and five specialists with 12 versioned skills on 1.47 million transaction rows. Passed 206/210 internal adversarial cases, including all 42 frozen cases.

Multimodal Content Creation Agent

Built evidence-based video repurposing with timestamps, quality gates, review, FFmpeg export, and recovery for product and educational content. CI includes real media tests.

SpatialMind

Independently built spatial memory, perception, and recovery. A three-seed, remembered-location 2D grid comparison reduced mean steps from 56.3 to 19.3 (−65.7%); broader baseline results are documented.

Financial LLM Fine-tuning

Built source-isolated financial data generation, two-stage LoRA / DeepSpeed SFT, numerical/factual evaluation, and dual-version vLLM serving. Published aggregate scores are documented with their evaluation conditions.

Reward Modeling Lab

Trained an 8B QLoRA reward model on one A10: frozen accuracy 50.42% to 91.35%, NDCG@5 0.947. Audited the stronger 94.61% length heuristic with controlled challenges.

Agentic-RL

Built a Text-to-SQL environment, answer isolation, Agent Lightning / veRL GRPO integration, and weight/budget controls. CPU checks pass; GPU policy-training results await validation.

Research

RewardLens · First author

Submitted to ICLR 2027 · Under review · [OpenReview ↗](#)

Proposed Behavioral Resolution, testing whether multimodal judges respond to relevant evidence through controlled image edits. Evaluated eight judges, four factors, 800 triplets, and 2,400 images. Five judges with perfect Spatial audit-base accuracy differed by 43.5pp in relevant-edit adaptation.

Education

Georgia Institute of Technology

M.S. in Computer Science · In progress

Aug 2026–Dec 2027 expected

Computational Perception & Robotics; interests in VLA, world models, and post-training.

Native Cantonese and Mandarin; professional working English, TOEFL 110.

Certifications: AWS Certified Solutions Architect – Professional · AWS Certified Machine Learning – Specialty

Nankai University

Bachelor's degree · International Economics and Trade

Sep 2014–Jun 2018

Public version updated: 2026-10-09