

Benjamin Taoson

本杰铭

AI engineering · Language models · Agent systems

benjamindaoson@gmail.com · github.com/Benjamindaoson

Profile

My experience spans enterprise integration, NLP algorithms, and language-model and agent engineering. I have worked on requirements, data processing, model design, tool and API integration, regression evaluation, and release validation. I focus on assigning work across small models, language models, and deterministic services according to task complexity, confidence, latency, and cost.

This public version covers technical responsibilities, skills, and public projects. A full résumé and further background are available by email.

Experience

Experience information updated September 2026.

Since 2023

AI algorithm engineering · LLM / Agent

- Worked on enterprise knowledge assistance, product questions, and product selection, covering requirements, data and algorithm design, agent architecture, API integration, release validation, and troubleshooting.
- Designed cooperation between small models, LLMs, agents, and deterministic services, with routing, retrieval, tool calling, task state, and context management.
- Used routing errors, low-confidence requests, and human corrections to develop evaluation cases, regression checks, failure analysis, and reusable components.

2021–2022

NLP algorithm engineering

- Built multilabel classification and information-extraction pipelines for business documents using MacBERT / BERT, CRF, hard-case feedback, and class-specific thresholds, with project-level train, validation, and test separation.
- Worked on semantic matching with bi-encoder retrieval, hard negatives, and cross-encoder reranking, alongside normalization of extracted amounts, units, dates, and comparison relations.

2018–2020

Digital solutions engineering

- Worked across sales, procurement, production, warehousing, and quality processes, contributing to requirements, workflow analysis, technical plans, system integration, and release acceptance.
- Integrated business-system data and APIs, using SQL and Python for data processing, problem analysis, and coordination across systems.

Technical skills

The links below provide public technical material relevant to these skills, presented separately from the enterprise experience above.

Models and training	Python, PyTorch, BERT-family models, LoRA / QLoRA, and DeepSpeed; classification, information extraction, preference modeling, and post-training. Reward Modeling Lab — Training and reward evaluation →
Agents and data systems	LangGraph, RAG, tool calling, task state and context management, SQL, PostgreSQL, API integration, and vLLM serving. Enterprise Data Agent — Semantic constraints and task orchestration →
Evaluation and verification	Data separation, preference and ranking evaluation, length-shortcut audits, controlled multimodal interventions, failure analysis, pytest, and API contract validation. RewardLens — Controlled visual interventions →

Public projects

[Enterprise Data Agent](#)

Connects business semantics, multi-agent investigation, deterministic analysis, and evidence-linked reports. The case study covers governed queries and traceable calculations, with the planning policy and data scope used in evaluation.

[Reward Modeling Lab](#)

QLoRA reward-model training and auditing, combining pairwise preferences, global ranking, and length-controlled challenges to investigate shortcuts. The repository distinguishes completed experiments from follow-up configurations that have not been run.

[RewardLens](#)

Holds questions and candidate answers fixed while editing visual evidence to examine adaptation and invariance in multimodal judges. The public manuscript, code, and result artifacts support findings within the recorded evaluation setting.

[View all work and research →](#)